

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: gpu=pd.read_csv("gpu_specs_v6.csv")
```

```
In [3]: gpu
```

Out[3]:

	manufacturer	productName	releaseYear	memSize	memBusWidth	gpuClock	memClock	un
0	NVIDIA	GeForce RTX 4050	2023.0	8.000	128.0	1925	2250.0	
1	Intel	Arc A350M	2022.0	4.000	64.0	300	1500.0	
2	Intel	Arc A370M	2022.0	4.000	64.0	300	1500.0	
3	Intel	Arc A380	2022.0	4.000	64.0	300	1500.0	
4	Intel	Arc A550M	2022.0	8.000	128.0	300	1500.0	
...
2884	3dfx	Voodoo5 5000 AGP	NaN	0.016	128.0	166	166.0	
2885	3dfx	Voodoo5 5000 PCI	NaN	0.016	128.0	166	166.0	
2886	3dfx	Voodoo5 6000	NaN	0.032	128.0	166	166.0	
2887	Intel	Xe DG1	NaN	4.000	128.0	900	2133.0	
2888	Intel	Xe DG1-SDV	NaN	8.000	128.0	900	2133.0	

2889 rows × 16 columns



In [4]: `gpu.head()`

Out[4]:

	manufacturer	productName	releaseYear	memSize	memBusWidth	gpuClock	memClock	unified
0	NVIDIA	GeForce RTX 4050	2023.0	8.0	128.0	1925	2250.0	
1	Intel	Arc A350M	2022.0	4.0	64.0	300	1500.0	
2	Intel	Arc A370M	2022.0	4.0	64.0	300	1500.0	
3	Intel	Arc A380	2022.0	4.0	64.0	300	1500.0	
4	Intel	Arc A550M	2022.0	8.0	128.0	300	1500.0	

In [5]: `gpu.shape`

Out[5]: (2889, 16)

In [6]: `gpu.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2889 entries, 0 to 2888
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   manufacturer          2889 non-null   object
 1   productName           2889 non-null   object
 2   releaseYear           2845 non-null   float64
 3   memSize               2477 non-null   float64
 4   memBusWidth           2477 non-null   float64
 5   gpuClock              2889 non-null   int64
 6   memClock              2477 non-null   float64
 7   unifiedShader         2065 non-null   float64
 8   tmu                   2889 non-null   int64
 9   rop                   2889 non-null   int64
10  pixelShader           824 non-null    float64
11  vertexShader          824 non-null    float64
12  igp                   2889 non-null   object
13  bus                   2889 non-null   object
14  memType               2889 non-null   object
15  gpuChip               2889 non-null   object
dtypes: float64(7), int64(3), object(6)
memory usage: 361.2+ KB
```

```
In [7]: gpu.rename(columns={"tmu":"Texture Mapping Unit"},inplace=True)
```

```
In [8]: gpu.rename(columns={"rop":"Render output unit","igp":"integrated graphics process
```

```
In [9]: gpu.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2889 entries, 0 to 2888
Data columns (total 16 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   manufacturer                          2889 non-null   object
 1   productName                          2889 non-null   object
 2   releaseYear                          2845 non-null   float64
 3   memSize                              2477 non-null   float64
 4   memBusWidth                          2477 non-null   float64
 5   gpuClock                             2889 non-null   int64
 6   memClock                             2477 non-null   float64
 7   unifiedShader                        2065 non-null   float64
 8   Texture Mapping Unit                 2889 non-null   int64
 9   Render output unit                   2889 non-null   int64
10   pixelShader                          824 non-null    float64
11   vertexShader                         824 non-null    float64
12   integrated graphics processor         2889 non-null   object
13   bus                                  2889 non-null   object
14   memType                              2889 non-null   object
15   gpuChip                              2889 non-null   object
dtypes: float64(7), int64(3), object(6)
memory usage: 361.2+ KB
```

```
In [10]: gpu.isnull().sum()
```

```
Out[10]: manufacturer          0
productName                    0
releaseYear                    44
memSize                        412
memBusWidth                    412
gpuClock                       0
memClock                       412
unifiedShader                  824
Texture Mapping Unit           0
Render output unit             0
pixelShader                    2065
vertexShader                   2065
integrated graphics processor   0
bus                             0
memType                         0
gpuChip                         0
dtype: int64
```

```
In [11]: gpu.drop(columns=["pixelShader","vertexShader"],inplace=True)
```

In [12]: `gpu.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2889 entries, 0 to 2888
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   manufacturer                        2889 non-null   object
 1   productName                        2889 non-null   object
 2   releaseYear                        2845 non-null   float64
 3   memSize                            2477 non-null   float64
 4   memBusWidth                        2477 non-null   float64
 5   gpuClock                           2889 non-null   int64
 6   memClock                           2477 non-null   float64
 7   unifiedShader                      2065 non-null   float64
 8   Texture Mapping Unit               2889 non-null   int64
 9   Render output unit                 2889 non-null   int64
10   integrated graphics processor      2889 non-null   object
11   bus                                2889 non-null   object
12   memType                            2889 non-null   object
13   gpuChip                            2889 non-null   object
dtypes: float64(5), int64(3), object(6)
memory usage: 316.1+ KB
```

In [13]: `gpu.isnull().sum()`

```
Out[13]: manufacturer      0
productName                0
releaseYear                44
memSize                    412
memBusWidth                412
gpuClock                   0
memClock                   412
unifiedShader              824
Texture Mapping Unit       0
Render output unit         0
integrated graphics processor 0
bus                        0
memType                    0
gpuChip                    0
dtype: int64
```

As memory size, memory bus width and memory clock specs are not available for the console based gpu and integrated gpu, therefore we will drop those rows because we only have to analyze the specifications of physical gpu units

In [14]: `gpu.dropna(inplace=True)`

```
In [15]: gpu.isnull().sum()
```

```
Out[15]: manufacturer      0
         productName      0
         releaseYear      0
         memSize          0
         memBusWidth      0
         gpuClock         0
         memClock         0
         unifiedShader    0
         Texture Mapping Unit  0
         Render output unit  0
         integrated graphics processor  0
         bus              0
         memType          0
         gpuChip          0
         dtype: int64
```

```
In [16]: gpu.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1721 entries, 0 to 2340
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   manufacturer                        1721 non-null   object
 1   productName                        1721 non-null   object
 2   releaseYear                        1721 non-null   float64
 3   memSize                            1721 non-null   float64
 4   memBusWidth                        1721 non-null   float64
 5   gpuClock                           1721 non-null   int64
 6   memClock                           1721 non-null   float64
 7   unifiedShader                      1721 non-null   float64
 8   Texture Mapping Unit                1721 non-null   int64
 9   Render output unit                  1721 non-null   int64
10   integrated graphics processor        1721 non-null   object
11   bus                                 1721 non-null   object
12   memType                            1721 non-null   object
13   gpuChip                            1721 non-null   object
dtypes: float64(5), int64(3), object(6)
memory usage: 201.7+ KB
```

```
In [17]: gpu["releaseYear"]=gpu["releaseYear"].astype("int64")
```

```
In [18]: gpu["memBusWidth"]=gpu["memBusWidth"].astype("int64")
```

```
In [19]: gpu["memClock"]=gpu["memClock"].astype("int64")
```

```
In [20]: gpu["unifiedShader"]=gpu["unifiedShader"].astype("int64")
```

In [21]: `gpu.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1721 entries, 0 to 2340
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   manufacturer                          1721 non-null   object
 1   productName                          1721 non-null   object
 2   releaseYear                          1721 non-null   int64
 3   memSize                              1721 non-null   float64
 4   memBusWidth                          1721 non-null   int64
 5   gpuClock                             1721 non-null   int64
 6   memClock                             1721 non-null   int64
 7   unifiedShader                        1721 non-null   int64
 8   Texture Mapping Unit                 1721 non-null   int64
 9   Render output unit                   1721 non-null   int64
10   integrated graphics processor        1721 non-null   object
11   bus                                  1721 non-null   object
12   memType                              1721 non-null   object
13   gpuChip                              1721 non-null   object
dtypes: float64(1), int64(7), object(6)
memory usage: 201.7+ KB
```

In [22]: `gpu.isnull().sum()`

```
Out[22]: manufacturer      0
productName              0
releaseYear               0
memSize                   0
memBusWidth               0
gpuClock                  0
memClock                  0
unifiedShader             0
Texture Mapping Unit      0
Render output unit        0
integrated graphics processor 0
bus                       0
memType                   0
gpuChip                   0
dtype: int64
```

In [23]: gpu

Out[23]:

	manufacturer	productName	releaseYear	memSize	memBusWidth	gpuClock	memClock	un
0	NVIDIA	GeForce RTX 4050	2023	8.000	128	1925	2250	
1	Intel	Arc A350M	2022	4.000	64	300	1500	
2	Intel	Arc A370M	2022	4.000	64	300	1500	
3	Intel	Arc A380	2022	4.000	64	300	1500	
4	Intel	Arc A550M	2022	8.000	128	300	1500	
...
2113	NVIDIA	Tesla S870	2007	1.536	384	600	800	
2114	ATI	Xbox 360 GPU 80nm	2007	0.512	128	500	700	
2157	NVIDIA	GeForce 8800 GTS 640	2006	0.640	320	513	792	
2158	NVIDIA	GeForce 8800 GTX	2006	0.768	384	576	900	
2340	ATI	Xbox 360 GPU 90nm	2005	0.512	128	500	700	

1721 rows × 14 columns



```
In [24]: gpu.describe()
```

Out[24]:

	releaseYear	memSize	memBusWidth	gpuClock	memClock	unifiedShader	Tex Map
count	1721.000000	1721.000000	1721.000000	1721.000000	1721.000000	1721.000000	1721.000000
mean	2013.571761	4.345259	322.584544	861.147008	1102.266124	1170.514817	71.453
std	4.130951	8.273479	761.177359	325.848376	407.848886	1769.078898	84.930
min	2005.000000	0.128000	32.000000	300.000000	266.000000	8.000000	4.000
25%	2010.000000	1.024000	128.000000	620.000000	800.000000	160.000000	20.000
50%	2013.000000	2.000000	128.000000	796.000000	1000.000000	480.000000	40.000
75%	2017.000000	4.000000	256.000000	1005.000000	1375.000000	1536.000000	96.000
max	2023.000000	128.000000	8192.000000	2331.000000	2257.000000	17408.000000	880.000



```
In [25]: #1.Which gpu were released in year 2013?
```



```
In [100]: a=gpu[(gpu["releaseYear"]>=2010)&(gpu["releaseYear"]<=2015)]
a
```

Out[100]:

	manufacturer	productName	releaseYear	memSize	memBusWidth	gpuClock	memClock	un
649	AMD	FirePro S9170	2015	32.000	512	930	1250	
650	AMD	FirePro W4130M	2015	1.024	128	775	1000	
651	AMD	FirePro W4150M	2015	1.024	128	800	1000	
652	AMD	FirePro W4170M	2015	2.000	128	825	1000	
653	AMD	FirePro W4190M	2015	2.000	128	825	1000	
...	
1673	AMD	Radeon HD 6850	2010	1.024	256	775	1000	
1674	AMD	Radeon HD 6870	2010	1.024	256	900	1050	
1675	AMD	Radeon HD 6950	2010	2.000	256	800	1250	
1676	AMD	Radeon HD 6970	2010	2.000	256	880	1375	
1678	ATI	Xbox 360 S GPU	2010	0.512	128	500	700	

866 rows × 14 columns




```
In [27]: #2.Which gpu has the highest memory clock?
```

```
In [38]: gpu[gpu["memClock"]==gpu["memClock"].max()]["productName"]
```

Out[38]: 437 GeForce GTX 1060 6 GB 9Gbps
Name: productName, dtype: object


```
In [29]: #3.Highest number of gpu released in a year?  
plt.figure(figsize=(9,9))  
gpu["releaseYear"].value_counts().plot(kind="bar")
```

Out[29]: <AxesSubplot: >




The highest number of gpu were released in year 2013 which is 182 which 2005 has the lowest release of only 1 gpu

```
In [101]: #4.Which manufacturer has more percentage of gpu models?  
y=gpu["manufacturer"].value_counts()  
plt.figure(figsize=(15,15))  
plt.pie(y,labels=["NVIDIA","AMD","ATI","INTEL"],colors=["green","red","yellow","orange"],  
plt.show()
```



Nvidia has more than 50% of market share of gpu which AMD has 34.81 percent of the share. while intel has lowest share of only 1.16%

```
In [36]: #5.Difference in memory clock speed with respect to year
plt.figure(figsize=(9,9))
plt.scatter(gpu["releaseYear"],gpu["memClock"])
plt.show()
```



As the technology advanced,gpu were provided with better and faster memory clock speeds which helps in faster rendering of graphics.

```
In [33]: #6.How many gpu has PCIe4.0x16 slots?
```

```
In [48]: a=gpu[gpu["bus"]=="PCIe 4.0 x16"]
a["bus"].value_counts()
```


```
Out[48]: PCIe 4.0 x16    106
Name: bus, dtype: int64
```

```
In [52]: #7. Most Common memory type  
gpu["memType"].value_counts().head(1)
```

```
Out[52]: GDDR5      712  
         Name: memType, dtype: int64
```


```
In [63]: #8.What is the most common memory size for PCIe 4.0x16?  
plt.figure(figsize=(7,7))  
a["memSize"].value_counts().plot(kind="bar")
```

```
Out[63]: <AxesSubplot: >
```




Most common memory size is 8gb for PCIe 4.0x16 slot which is 24

```
In [93]: #9. Which gpu in year 2020 has the highestTexture Mapping Unit?
a=gpu[(gpu["releaseYear"]==2022)]
plt.figure(figsize=(12,12))
plt.barh(a["productName"],a["Texture Mapping Unit"],color="orange")
plt.show()
```



As we can see from the above bar plot Radeon RX 7900XT has the highest texture mapping unit released in the year 2020 while the lowest are steam deck gpu and Geforce MX550

```
In [120]: #10. Highest memory size in year 2022
m=gpu[gpu["releaseYear"]==2022]
y=m["memSize"]
x=m["productName"]
plt.figure(figsize=(15,12))
plt.barh(x,y,label="gpu CLock")
plt.show()
```



As we can see highest memory size is 80gb which are H100SXM5 and H100PCIe. These gpu are only available for enterprise and data server market and are not available for common consumer. Considering the gpus available for common consumer highest memory size in 2022 is 24gb while lowest is 2gb

```
In [124]: #11. Maximum memory size
gpu["memSize"].max()
```

Out[124]: 128.0